

Email: ganapathysubramaniam1999@gmail.com
Phone #: +1-437-556-0264
Location: Toronto

Checkout my recent Gen AI personal Featured section of my [linkedin](#) and my [website](#)

EXPERIENCE

AI/ML Development Intern | Nokia | Canada | January 2025 – Present

- **Designed AI Architectures:** Engineered cutting-edge software architectures for in-house AI products using **RAG**, ensuring seamless knowledge retrieval and peak system performance.
- **Built LLMops Pipelines:** Created and optimized scalable LLM-Ops pipelines, streamlining model lifecycle from fine-tuning to deployment and monitoring.
- **Optimized LLM Performance:** Monitored and enhanced LLM metrics, driving accuracy, efficiency, and real-time optimization.
- **Developed Vector Database Schemas:** Crafted advanced vector database schemas, enabling superior retrieval and scalability for AI-driven solutions.
- **Leveraged Azure ML Studio** to design and manage impactful AI workflows and models for business transformation.
- **Created Autonomous AI Agents:** Designed intelligent AI agents capable of automating tasks, making decisions, and driving operational efficiency across diverse applications.
- **Innovated AI Products:** Delivered proprietary AI solutions blending advanced technologies with user-focused design to solve complex challenges.

AI/ML Developer | Credwise | Canada | Volunteering | August 2024 - Present

- **Led a high-performing team of 6 AI/ML engineers**, delivering 3 major Generative AI initiatives that achieved 40% faster data processing and 25% reduced response times in AI-powered applications.
- Architected and implemented **advanced Retrieval-Augmented Generation (RAG)** systems integrated with Large Language Models on **Vertex AI**, enhancing real-time inferencing and model precision.
- Engineered scalable cloud infrastructure using GCP services, **optimizing Cloud Firestore databases for real-time synchronization**, automated scaling, and low-latency data retrieval under high-traffic conditions.
- Established **containerized AI model deployments** using Docker and Kubernetes on Google Cloud Platform, implementing efficient CI/CD pipelines for continuous model updates and system monitoring.
- Designed and **optimized cloud-native AI architectures** integrating multiple data streams, focusing on performance tuning, automated scaling, and robust application development.

Python Engineer | Virtusa | IN, Tamil Nadu, Chennai | October 2021 - July 2023

- Engineered high-performance Python frameworks that achieved 20% improvement in application performance through algorithmic optimizations and code refactoring, **including development of a cost-effective ETL framework that replaced Talend for 170% cost reduction**.
- Architected AWS cloud infrastructure solutions using Python and boto3, implementing automated resource provisioning systems and **time series forecasting models** for proactive cost management and optimization of cloud resources.
- Developed comprehensive data warehousing solutions on **Snowflake with Python-driven ETL processes**, resulting in 20% faster data processing and enhanced data-driven decision-making capabilities.
- Created **automated data pipeline** systems that increased processing efficiency by 30%, incorporating robust data cleaning, transformation, and validation processes while ensuring data integrity across multiple stages.
- Designed and implemented **real-time web frameworks** using Python and Flask, enabling dynamic system configuration and enhanced user interaction capabilities for improved customer engagement.

Junior Software Developer | Powertrac Engineers Pvt Ltd | IN, Tamil Nadu, Chennai | May 2021 - September 2021

- Successfully **deployed 5 enterprise software applications** through comprehensive end-to-end development cycles, achieving 20% improvement in operational efficiency through collaborative development practices.
- Administered and **optimized multiple MySQL databases**, implementing performance tuning strategies and maintaining high availability standards for business-critical operations.
- Developed robust software solutions using Java and Python, focusing on functionality enhancement and system optimization for improved application performance.
- Implemented continuous integration practices and automated testing frameworks, resulting in **accelerated development cycles** and significant reduction in defect rates.
- Led cross-functional collaboration initiatives for software improvement projects, contributing to enhanced team productivity and streamlined development processes

Key Strengths

- **Architecting End-to-End LLM Solutions:** Expertise in delivering comprehensive LLM solutions across diverse use cases.
- **Optimizing LLM Performance:** Proficient in advanced tuning and integration techniques to enhance model efficiency.
- **Leveraging Cutting-Edge AI:** Skilled in applying the latest AI technologies to solve complex business challenges.
- **Rapid Adaptation to Emerging Technologies:** Demonstrated ability to quickly adapt and excel in new AI frameworks and cloud services.
- **Sophisticated LLM Applications:** Proficient in using LangChain and LangGraph to develop advanced LLM applications.

Soft Skills & Attributes

- **Problem-Solving:** Strong analytical thinker with exceptional debugging skills
- **Collaboration:** Thrives in team environments, bringing creative solutions to complex challenges
- **Self-Learner:** Eager to embrace new technologies and methodologies with a growth mindset
- **Innovative & Proactive:** Passionate about using technology to create a positive global impact

| Category | Skill Area | Specific Technologies & Expertise |
|--|-----------------------------------|---|
| Large Language Models (LLMs) & AI Technologies | LLM Development & Deployment | <ul style="list-style-type: none"> • Fine-Tuning Techniques: - RAFT - PEFT - Traditional approaches- RAG-Based Architecture: - Utilizing vector databases for retrieval-augmented generation- Cloud-Based Deployment: - Expertise in deploying models via Vertex AI - AWS Sagemaker endpoints • Hyperparameter Optimization: - Advanced techniques for both training and inference • Prompt Engineering: - Crafting prompts for structured and coherent response generation |
| | LLM Frameworks & SDKs | <ul style="list-style-type: none"> • Open-Source: - Hugging Face Transformers – PyTorch • Proprietary: : OpenAI (GPT), Anthropic (Claude), Google (Gemini) • LLM Application Frameworks: - LangChain and LangGraph • Agent-Based Systems: - CrewAI |
| | Multimodal AI | <ul style="list-style-type: none"> • Image Generation: - Stability AI - OpenAI DALL-E 3 • Speech-to-Text: - Faster-Whisper - OpenAI Whisper – Google STT • Text-to-Speech: - OpenAI TTS-HD, - Google Cloud TTS – Deepgram • Translation: - Google Translation API • Voice AI Agent: Openai Realtime, Twilio |
| Data Science & Machine Learning | Machine Learning Algorithms | <ul style="list-style-type: none"> • Regression Algorithms- Classification Algorithms- Clustering Algorithms using Sci-kit Learn |
| | Time Series Forecasting | <ul style="list-style-type: none"> • ARIMA- SARIMA- SARIMAX |
| | Natural Language Processing (NLP) | <ul style="list-style-type: none"> • NLTK- Textblob- VADER- Gensim- Textacy |
| | Data Reporting & Analytics | <ul style="list-style-type: none"> • Power BI |
| Programming & Data Technologies | Programming Languages | <ul style="list-style-type: none"> • Primary: - Python |
| | Databases | <ul style="list-style-type: none"> • NoSQL: - MongoDB - Google Firestore • Relational: - SQLite3 - MySQL |
| | Cloud Data Warehousing | <ul style="list-style-type: none"> • Snowflake |
| Cloud Platforms & Integration | Cloud Expertise | <ul style="list-style-type: none"> • Google Cloud Platform (GCP): - Vertex AI Agents - Vector Search - Cloud Run - Storage Buckets – Firestore • Amazon Web Services (AWS): - AWS Bedrock - Sagemaker - S3 buckets - IAM - Lambda – DynamoDB • Microsoft Azure: Azure ML Studio – Model catalog -Azure Openai Services • Other AI Platforms: - Gemini’s AI Studio - OpenAI Console- Anthropic developer console- Github Models |

EDUCATION

Postgraduate Degree, Artificial Intelligence and Machine Learning | Lambton College | CA, ON, Toronto | 2025

- Lambton College, Toronto | Expected August 2025
- *in The Dean's Honor List in Semester 1 and Semester 2*

Bachelor of Engineering (BE), Computer Science | St. Joseph's College of Engineering | 2020

- St. Joseph's College of Engineering, India | July 2020 (Due to the pandemic in 2020, the government rescheduled the actual 8th semester exam to November 2020, ensuring continuity in academic assessments.)

CERTIFICATIONS

| Course | Organization | |
|--|--------------|---|
| Advanced Python | StudySection | https://www.studysection.com/users/socialMedia/NzMzMjI4/Z2FuYXBhdGh5c3VicmFrYW5pYW0xOTk5QGdtYWlsLmNvbQ== |
| Python | TestDome | www.testdome.com/certificates/1673961486f8444594a87039fe1314c0 |
| Python | LinkedIn | Skill Assessment Badge |
| Machine Learning | LinkedIn | Skill Assessment Badge |
| Mysql | LinkedIn | Skill Assessment Badge |
| Data Ops Fundamentals | DATA KITCHEN | https://www.linkedin.com/posts/ganapathy-subramaniam-sundar-b08aa222b_dataops-certification-activity-6892886657306017793-_lmp?utm_source=share&utm_medium=member_desktop |
| Dataops platform developer certification | DATA KITCHEN | https://www.linkedin.com/posts/ganapathy-subramaniam-sundar-b08aa222b_dataops-platform-developer-certification-activity-6893635485772644352-ZgSu?utm_source=share&utm_medium=member_desktop |
| Problem Solving using Python | HackerRank | https://www.hackerrank.com/certificates/cb38b70513ac |
| Python | HackerRank | https://www.hackerrank.com/certificates/376aba178b80 |
| Snowflake Data warehouse | Udemy | https://www.linkedin.com/posts/ganapathy-subramaniam-sundar-b08aa222b_snowflake-activity-6893935118860247040-mhNJ?utm_source=share&utm_medium=member_desktop |